

UNIVERSITY OF PENNSYLVANIA

ESE 546: PRINCIPLES OF DEEP LEARNING

FALL 2020

[11/30] HOMEWORK 5

DUE: 12/15 TUE 11.59P ET

Changelog

Instructions

Read the following instructions carefully before beginning to work on the homework.

- You will submit solutions typeset in L^AT_EX on Gradescope (strongly encouraged). You can use hw_template.tex on Canvas in the “Homeworks” folder to do so. If your handwriting is *unambiguously legible*, you can submit PDF scans/tablet-created PDFs.
- Please start a new problem on a fresh page and mark all the pages corresponding to each problem. Failure to do so may result in your work not graded completely.
- Clearly indicate the name and Penn email ID of all your collaborators on your submitted solutions.
- For each problem in the homework, you should mention the total amount of time you spent on it.
- You can be informal while typesetting the solutions, e.g., if you want to draw a picture feel free to draw it on paper clearly, click a picture and include it in your solution. Do not spend undue time on typesetting solutions.
- You will see an entry of the form “HW 5 PDF” where you will upload the PDF of your solutions. You will also see an entry named “HW 5 Problem 2 Code” where you will upload the code for Problem 2. **For each programming problem, you should create a fresh Google Colab notebook.** This notebook should contain all the code to reproduce the results of the problem and you will upload the .ipynb file obtained from Colab. Name your notebook to be “pennkey_hw5_problem4.ipynb”, e.g., I will name my code for Problem 2 as “pratikac_hw5_problem2.ipynb”.
- **Remember that you should include all the relevant plots in the PDF, without doing so you will not get full credit. Your PDF solutions should be completely self-contained, we should not have have to look at/run the Colab notebook to check your solutions.**

Credit The points for the problems add up to 100. You only need to solve for 50 points to get full credit, i.e., your final score will be $\min(\text{your total points}, 50)$.

- 1 **Problem 1 (40 points, background on variational inference).** Given some distribution

$$p(w) = \frac{e^{-\beta\Phi(w)}}{Z}$$

- 2 variational inference approximates $p(w)$ using some distribution of our choice, denoted by $q(w)$, by
3 solving the problem

$$\min_{q \in \mathcal{Q} \subset \mathcal{P}(\mathcal{W})} \text{KL}(q \parallel p). \quad (1)$$

- 4 The set \mathcal{Q} is called the variational family which is typically a subset of all probability distributions
5 $\mathcal{P}(\mathcal{W})$ on the domain $\mathcal{W} \ni w$. We should take \mathcal{Q} to be small enough that we can easily solve the
6 optimization problem (1) but still large enough to obtain a small $\text{KL}(q \parallel p)$.

- 7 The quantity

$$\mathcal{F}(q) := -\beta \sum_{w \in \mathcal{W}} q(w) \Phi(w) - \sum_{w \in \mathcal{W}} q(w) \log q(w)$$

- 8 is called the free-energy by physicists. For distributions p, q supported on some discrete set, we can
9 rewrite $\text{KL}(q \parallel p)$ as

$$\text{KL}(q \parallel p) = \beta \sum_{w \in \mathcal{W}} q(w) \Phi(w) + \sum_{w \in \mathcal{W}} q(w) \log q(w) + \log Z(\beta). \quad (2)$$

- 10 Since $\text{KL}(q \parallel p) \geq 0$ we immediately have a lower bound for $\log Z(\beta)$

$$\max_{q \in \mathcal{Q}} \mathcal{F}(q) \leq \log Z(\beta) = \log \sum_{w \in \mathcal{W}} e^{-\beta\Phi(w)}. \quad (3)$$

- 11 This is a very useful bound because the log-partition function $\log Z(\beta)$ contains essentially all
12 information about the probability distribution $p(w)$. For example, the average value of $\Phi(w)$ is given
13 by

$$\mathbb{E}_{w \sim p(w)} [\Phi(w)] = -\frac{\partial \log Z(\beta)}{\partial \beta} \approx \mathbb{E}_{w \sim q(w)} [\Phi(w)]. \quad (4)$$

- 14 The last approximation is reasonable because $q(w) \approx p(w)$ if we have done a good job of minimizing
15 $\text{KL}(q \parallel p)$. If $\mathcal{Q} = \mathcal{P}(\mathcal{W})$, we certainly have

$$\max_{q \in \mathcal{P}(\mathcal{W})} \mathcal{F}(q) = \log Z(\beta).$$

- 16 because in this case the minimum of $\text{KL}(q \parallel p)$ is zero.

Note: You will notice that the bound on the log-partition function is very similar to the bound on $p(x)$ we obtained when we derived ELBO (x denotes data and z denotes the representation in the equation)

$$\text{ELBO}(q) = \mathbb{E}_{z \sim q(z|x)} [\log p(x|z)] - \text{KL}(q(z|x) \parallel p(z)) \leq \log p(x).$$

This is not be surprising because $p(x)$ is indeed the normalizing constant in Bayes law

$$p(z|x) = \frac{p(x,z)}{p(x)} = \frac{p(x,z)}{\sum_z p(x,z)}.$$

- 17 (a) **(5 points)** Prove the expression in (2).

- 18 (b) **(5 points)** Prove the equality in (4).

19 (c) **(10 points)** The mean-field assumption in variational inference means that we pick the family of
 20 distributions that factors over the variable w . One example is,

$$\mathcal{Q}_1 = \left\{ q(w) : q(w) = \prod_{i=1}^N q_i(w_i) \right\}.$$

21 where $q_i(w_i)$ is a distribution which only depends on the i^{th} variable w_i . The distributions q_i can also
 22 be the same in which case we would have

$$\mathcal{Q}_2 = \left\{ q(w) : q(w) = \prod_{i=1}^N q(w_i) \right\}.$$

23 As you can imagine, it is difficult to pick a variational family without knowing much about the
 24 distribution $p(w)$. This problem will consider a simple example. Let true data be generated by

$$p(w) \propto 1 - \prod_{i=1}^N w_i$$

25 where $w_i \in \{0, 1\}$. Clearly the configuration $w_i = 1$ for all $i = 1, \dots, N$ has zero probability, all
 26 other configurations have equal probability $1/(2^N - 1)$. This would seem like a simple distribution
 27 $p(w)$ but the highly correlated part (one configuration) makes it difficult to approximate using
 28 mean-field families. Let's pick \mathcal{Q}_2 as the family. Find the minimizer of

$$q^* = \min_{q \in \mathcal{Q}_2} \text{KL}(q \parallel p).$$

29 Does $q^*(w)$ look similar to $p(w)$?

30 (d) **(20 points)** Now pick a different mean-field family \mathcal{Q}_1 and solve

$$q^* = \min_{q \in \mathcal{Q}_1} \text{KL}(q \parallel p).$$

31 It will help to notice that $q(1, 1, \dots, 1) = 0$ is a necessary condition for the minimizer (why?). Notice
 32 that the entire problem is symmetric in all the variables. So one way of ensuring our solution satisfies
 33 the necessary condition is to simply set $q_1(w_1) = 1 - w_1$. Now does $q^*(w)$ look similar to $p(w)$?

34 **Problem 2 (60 points. Variational Auto-Encoder for MNIST. Do this on Colab/AWS but this
 35 can also be done on your laptop.).** In this problem we will train a variational auto-encoder for
 36 generating MNIST digits. The setup is exactly the same as that of Section 3 in the assigned reading
 37 “Auto-Encoding Variational Bayes” (<https://arxiv.org/abs/1312.6114>) by Kingma & Welling.

38 (i) **(0 points)** Create the dataset. Use 1000 images of each MNIST digit; this gives a total of 10,000
 39 images. Sub-sample the MNIST images to 14×14 size. You will also binarize the MNIST
 40 dataset for this problem: if a pixel is greater than 128 set it to 1, else set it to zero.

41 (ii) **(40 points)** The encoder will consist of two fully-connected layers, the first has $14 \times 14 = 196$
 42 inputs and 128 outputs (and tanh nonlinearity) and the second layer has 128 inputs and 16
 43 outputs (and no nonlinearity). Therefore, the latent factor $z \in \mathbb{R}^8$ (8 output neurons for the
 44 mean and 8 more for the standard deviation). The decoder takes in as input $z \in \mathbb{R}^8$, pushes
 45 it through one layer with 128 outputs (and tanh nonlinearity) and then another layer with 196
 46 output neurons (with sigmoid nonlinearity). Note that the final nonlinearity of the decoder
 47 should be sigmoid because we want to reconstruct binarized MNIST images. See Sections C1
 48 & C2 in Appendix of the above paper for more details.

49 Let the parameters of the encoder be u and the parameters of the decoder be v . We will
50 maximize the objective

$$\frac{1}{n} \sum_{i=1}^n \ell(u, v; x^i)$$

51 for images $\{x^1, \dots, x^n\}$ with respect to its parameters u, v where

$$\ell(u, v; x) = \mathbb{E}_{z \sim p_u(z|x)} [\log p_v(x|z)] - \text{KL}(p_u(z|x) \parallel N(0, I)). \quad (5)$$

52 and x is one MNIST datum. We will set

$$p_u(z|x) = N(\mu_u(x); \sigma_u(x)^2 I)$$

53 where $[\mu_u(x), \sigma_u(x)] \in \mathbb{R}^{16}$ is the output of encoder upon feeding the image x . We are
54 modeling $p_u(z|x)$ as a Gaussian distribution and the neural network predicts the mean $\mu_u(x)$
55 of the distribution and the diagonal of the covariance $\sigma_u(x)$. The KL-divergence is

$$\text{KL}(p_u(z|x) \parallel N(0, I)) = -\frac{1}{2} \sum_{i=1}^8 (1 + \log(\sigma_u(x)_i^2) - \mu_u(x)_i^2 - \sigma_u(x)_i^2).$$

56 We will use a Bernoulli distribution to model $p_v(x|z)$ because we are using the binarized
57 MNIST dataset:

$$\log p_v(x|z) = \sum_{i=1}^{196} x_i \log y_i + (1 - x_i) \log(1 - y_i)$$

58 where $\mathbb{R}^{196} \ni y = \text{output of the decoder}$.

59 Train the auto-encoder using the re-parametrization trick for computing the gradient of the u
60 and standard back-prop for computing the gradient of the $\log p_v(x|z)$. You should use 2 samples
61 to compute the expectation over $z \sim p_u(z|x)$.

- 62 (iii) **(5 points)** Plot the first and second term of ELBO in (5) separately as a function of the number
63 of weight updates.
- 64 (iv) **(5 points)** Pick 8 MNIST images, run them through the encoder and the decoder to plot the
65 output of the decoder side-by-side with the original images.
- 66 (v) **(5 points)** Sample from the generative model directly: sample $z \in \mathbb{R}^8$ from a standard Gaussian
67 distribution and run the decoder network for this z to synthesize an MNIST image. Plot the
68 synthesized images for a few different samples of z ; remember that you can also use a mini-batch
69 of latent variables z to synthesize a few images in one forward-prop of the decoder.
- 70 (vi) **(5 points)** We will next compute the validation performance of a generative model. Compute the
71 average validation log-likelihood $\log p(x|z)$ for 100 synthesized images after every 100 weight
72 updates. Draw a plot of the reconstruction log-likelihood term in ELBO on the training mini-
73 batches (one point after every weight update) and compare it with the validation log-likelihood
74 as training progresses (one point after every 100 weight updates).